

Sitecore Agentic Studio – Designing Scalable Agentic UIs for Enterprise AI

Role: Product Designer

Timeline: August 2025 – January 2026

Product: Sitecore Agentic Studio (within SitecoreAI, formerly XM Cloud)

Executive Summary

Sitecore Agentic Studio was the productisation of rapid, experimental AI work emerging from Sitecore AI Innovation Labs into a scalable, enterprise-ready platform. Its goal was to enable AI agents, flows, and a universal chat experience that could be embedded across SitecoreAI without slowing down complex marketer workflows or locking the product into brittle UI patterns.

As the Product Designer, I was responsible for translating fast-moving, demo-driven Innovation Labs concepts into production-ready experiences that could scale across unknown future agentic use cases. The core challenge was not just visual design, but systems design: creating UI structures that could support non-linear workflows, deep customisation, collaboration, and governance while still offering the familiarity and speed of chat.

This case study focuses on three intertwined problems:

- Designing scalable, highly customisable UIs for agents and flows
 - Integrating chat as a first-class but non-dominant interaction model
 - Bridging experimentation and enterprise readiness through cross-team communication
-

Context & Problem Space

Company & Product Context

Agentic Studio was the bridge between **Sitecore AI Innovation Labs** and **SitecoreAI**, forming the foundation for how AI would be delivered across the product suite. Strategically, this work supported Sitecore's shift from a composable product portfolio to a consolidated AI-first platform, in response to increasing competitive pressure.

Agentic Studio introduced:

- AI agents and flows that could be configured and reused
- A universal chat experience accessible from anywhere in SitecoreAI
- Purpose-built UIs for agentic workflows beyond chat

Agents and flows could be triggered from dashboards, managed centrally in the Agentic area, or interacted with through chat, creating a uniquely complex set of touchpoints.

Why This Was Not a “Normal” Design Project

This was not a single feature, but a new product surface being embedded into an existing enterprise ecosystem. Key differences from a typical product design project included:

- Extreme ambiguity due to the pace of AI innovation
- Constant inflow of new requirements from customers and industry changes
- Tight timelines (≈ 3 months to early customer access)
- Direct collaboration with a VP who was also the primary engineer

Design decisions had long-term architectural implications, often affecting how all future AI features would be built.

Role & Responsibilities

My Role in Practice

Officially, my role was **Product Designer**. In practice, I:

- Designed all Agentic Studio and chat experiences in Figma or prototyped them using code
- Led interaction and systems thinking for agents, flows, and chat
- Conducted competitor and industry research (e.g. Opal, Google and Microsoft AI tooling)
- Acted as the design bridge between Innovation Labs and SitecoreAI
- Owned handover quality, design sign-off, and implementation reviews

I worked most closely with:

- **VP of AI & Innovation** – requirements, customer input, rapid iteration
- **Design Lead (Agentic)** – design quality, patterns, and direction
- **Product Design Manager** – cross-product alignment and impact

Decision-making was highly collaborative but fast. I often had to propose a direction based on previous knowledge from Innovation Labs, test it quickly through design, and use that artefact to drive alignment.

Core UX Problems

Rather than framing this as a list of features, the work revolved around several fundamental UX problems.

1. How Do You Represent Agents and Flows?

Early on, we struggled with a binary choice: chat or custom UI. Chat offered speed and flexibility, but broke down for complex, non-linear workflows. Custom UIs offered clarity, but risked being brittle and expensive to maintain.

The key insight was that we did not need to choose.

By designing from **small UI modules** mapped to agent and flow configurations, we could:

- Render the same agent in chat or full UI
- Avoid duplicating logic
- Support future, unknown configurations without rebuilding UIs

2. How Do You Handle Non-Linear Workflows?

Unlike chat, agent workflows are not inherently linear. Users could:

- Skip steps
- Return later
- Collaborate asynchronously

A purely chat-based model created temporal problems and broke mental models. We addressed this by introducing **Spaces**, an abstraction layer where:

- Chat remained linear and familiar
- Full agent UIs could be non-linear
- Collaboration, reviews, and history lived outside chat

This preserved the simplicity of chat while enabling enterprise workflows.

3. How Do You Represent Enterprise Content?

Sitecore users work with structured content types (sites, collections, media), not just text. Traditional AI chats do not handle this well.

We introduced **Artifacts**: a flexible abstraction capable of representing any generated content type with tailored editing experiences. Artifacts could be grouped, reviewed, and approved consistently, then promoted into SitecoreAI as first-class content.

Why Traditional Patterns Failed

Standard workflow and chat patterns broke down under:

- Non-linear temporality
- Role-based permissions and visibility
- Parallel agent execution
- High output volume (hundreds of artifacts)

Relying solely on chat would have created an opaque, fragile system that slowed users down. The solution required combining chat familiarity with explicit, purpose-built interfaces.

Design Principles

Several principles guided decisions throughout the project:

- **Design for UI first, translate to chat later**
Optimal workflows should not be constrained by chat metaphors.
- **Atomic over bespoke**
Components must scale combinatorially without UI rewrites.
- **Linear chat, non-linear spaces**
Preserve mental models while enabling complexity.
- **Progressive disclosure**
Reduce cognitive load when dealing with high output volume.
- **Abstraction where stable, explicit control where volatile**
(e.g. Spaces vs Agent Builder).

Designing for Scalability & Customisation

Early Mental Models and Why They Changed

Initially, agents and flows were treated as the same concept. This broke down once we introduced collaboration and handovers. A flow could orchestrate agents; an agent could not

contain other agents. While technical, this distinction had real UX consequences, especially in builder experiences.

Managing Customisation Risk

Every customisation option multiplied design and technical debt. To control this, we:

- Designed from atomic components
- Created shared status signifiers and states
- Avoided one-off UI variants

This allowed us to support complex configurations without exponential complexity.

Delivering the Chat Experience

Why Chat Was Essential

Chat was not a nice-to-have. It provided:

- A safe baseline for unknown future use cases
- A way for users to add rich context beyond forms
- Familiarity in a fast-moving AI landscape

Embedding Chat into an Enterprise Product

We designed:

- A **universal chat**, accessible via the top navigation
- A **contextual prompt bar**, scoped to Agentic work

Key decisions included:

- Push-in side panel vs overlay
- Undocking chat on builder pages
- Redirecting actionable work into Spaces for review

This prevented chat from becoming a black box while preserving traceability and collaboration.

Cross-Team Communication & Influence

A major part of my role was translating between two modes of working:

- **Innovation Labs:** rapid, demo-driven, minimal constraints
- **Core Product:** governed, consistent, enterprise-ready

To align teams, I relied heavily on:

- Interactive prototypes for narrative alignment
- Annotated handovers with walkthrough videos
- Follow-up sessions to surface ambiguity early

Pushback—especially from product leadership—was common and healthy. Decisions around chat and agent interaction patterns had platform-wide implications, requiring clear rationale and iteration.

Outcomes & Impact

Product Impact

- Established Agentic Studio as the foundation for AI across SitecoreAI
- Enabled customers to build and scale their own AI agents
- Early partners built 100+ agents within weeks of release

Organisational Impact

- Centralised AI ownership within Agentic
- Reduced duplicated AI efforts across product teams
- Created a reusable system for future AI features

Trade-offs Accepted

- Prioritised UI-based workflows over refining chat for initial release
 - Accepted that chat required further iteration post-launch
-

Reflection

This project reinforced several lessons:

- Designing AI systems means designing for uncertainty
- Technical constraints should follow user flows—not the reverse

- Simplicity is harder than complexity, especially in AI

If starting again, I would invest earlier in identifying the most common agent use cases to ground atomic design decisions in concrete scenarios.

Why This Case Study Matters

This project demonstrates my ability to:

- Design scalable systems under extreme ambiguity
- Translate experimental AI concepts into enterprise UX
- Influence cross-functional teams through design thinking

It actively challenges the misconception that AI UX is just chat—and shows how thoughtful systems design can unlock AI's real value in complex products.